

# Running Django apps on Docker Cloud

Felix Kerekes - REBS



## About me - Felix Kerekes

- Coding in Python for ~8 years
- Started using Docker around 1 year ago

Google



rently.ro

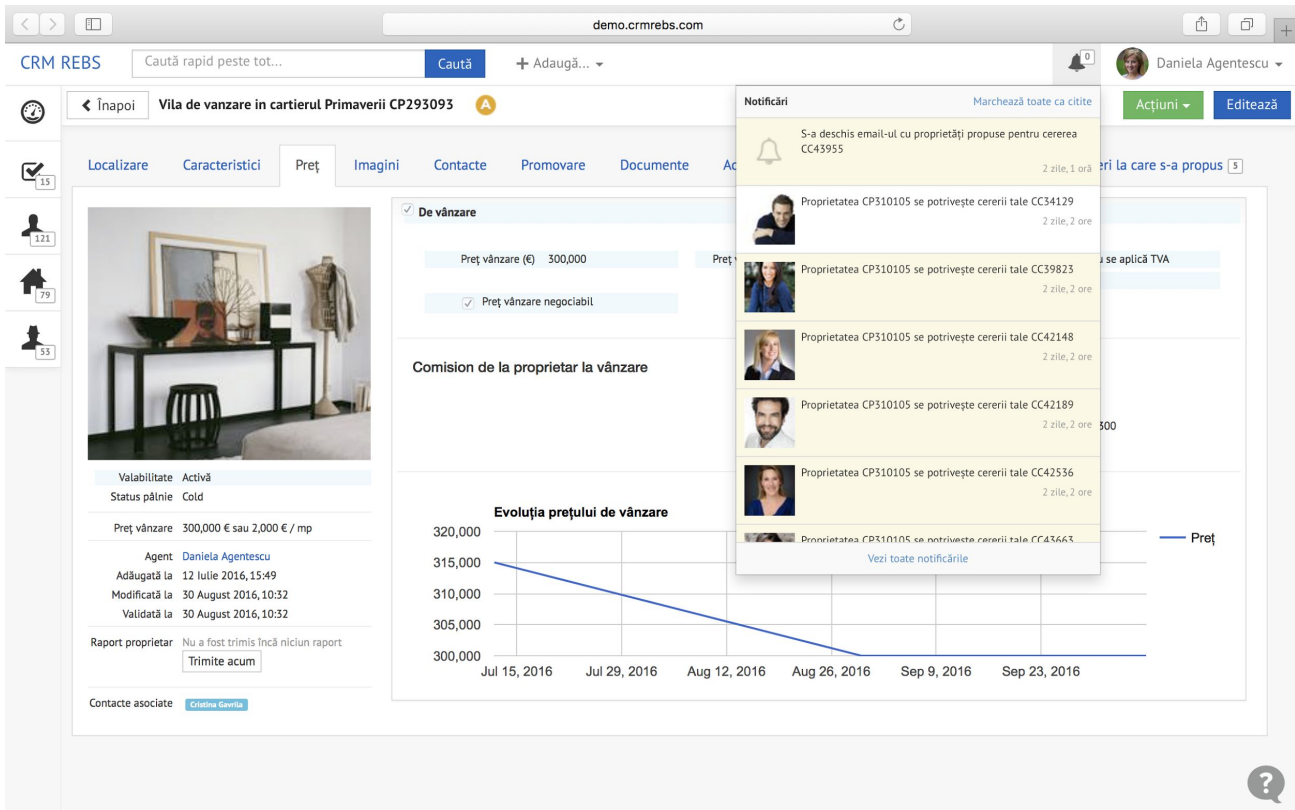


# REBS

Real Estate Business Solutions



# The product - CRM REBS





# Early days - Fabric for a single host

```
from fabric.api import *
from fabric.contrib import django
```

```
django.project('crmrebs')
```

```
env.hosts = ["YOUR_HOST_HERE"]
```

```
def deploy(name, tag='master'):
    set_env(name)
```

```
    with cd(env.path):
        run('git pull')
        run('git checkout %s' % tag)
```

```
    run('%s/bin/pip install -r %s' % (env.venv_path, env.requirements))
```

```
    with cd(env.path):
        run_in_virtualenv('manage.py collectstatic --noinput')
        run_in_virtualenv('manage.py compilemessages')
        run_in_virtualenv('manage.py compress -e html -e js --force')
        run_in_virtualenv('manage.py syncdb')
        run_in_virtualenv('manage.py migrate --merge')

    sudo('service crm-%s restart' % name)
    sudo('service crm-queue restart')
    sudo('service crm-queue-priority restart')
```



# Going multi-host



SALTSTACK

crmrebs-app:

git.latest:

- name: {{ pillar['git']['url'] }}
- user: {{ pillar['uwsgi']['user'] }}
- target: {{ pillar['django']['path'] }}
- identity: {{ pillar['git']['deployment\_key'] }}
- rev: {{ pillar['git']['rev'] }}
- force\_reset: True
- require:
  - pkg: version-control-pkgs
  - file: deployment-key

crmrebs-app-virtualenv:

virtualenv.managed:

- name: {{ pillar['django']['virtualenv'] }}
- user: {{ pillar['uwsgi']['user'] }}
- requirements: {{ pillar['django']['requirements'] }}
- user: {{ pillar['uwsgi']['user'] }}
- pip\_download\_cache: {{ pillar['django']['virtualenv'] }}.cache
- cwd: {{ pillar['django']['path'] }}
- watch:
  - git: crmrebs-app
- require:
  - pkg: crmrebs-app-virtualenv-dependencies
  - git: crmrebs-app
  - sls: python



# Going multi-host



**SALTSTACK**

- First time we could actually deploy configuration & infrastructure changes
- ~ 4000 lines of configuration
  - Managing the whole stack (PostgreSQL, Memcached, Redis, RabbitMQ, uWSGi, nginx)
  - Prod and staging environment
  - Other self-hosted tools (Phabricator, Jenkins, Sentry etc.)
- Harder to scale
- Personally, it's ugly
  - A lot of templated configuration files



# The problems we faced

- Scaling
- Redundancy
- Flexibility





## The solution - Docker





# The solution - Docker

```
FROM buildpack-deps:xenial
```

```
MAINTAINER Felix Kerekes <felix@crmrebs.com>
```

```
# Install dependencies
```

```
RUN apt-get update && apt-get install -y --no-install-recommends \
```

```
...
```

```
python-pip \
```

```
sudo \
```

```
&& apt-get clean \
```

```
&& rm -rf /var/lib/apt/lists/* /tmp/* /var/tmp/*
```

```
# Upgrade pip and install wheel
```

```
RUN pip install -U pip
```

```
RUN pip install wheel
```

```
# Install pip requirements
```

```
COPY requirements.txt /app/
```

```
RUN pip install --no-cache-dir -r requirements.txt
```

```
# Add application code
```

```
COPY . /app/
```

```
# Compile translation files
```

```
RUN ./manage.py compilemessages
```

```
EXPOSE 8000
```

```
CMD ./docker-entrypoint.sh
```



# The solution - Docker Cloud



- Formerly Tutum, acquired by Docker
- CaaS - Containers as a Service
- Orchestration - multiple deployment strategies (emptiest node, high availability, every node)
- Multi-host networking
- Volumes
- Terminology
  - Node - tags
  - Stack - Stackfile (fig / docker-compose like)
  - Service
  - Container



# Docker Cloud - Stackfiles



crmrebs-prod-memcached:

```
command: memcached -m 8192
image: 'memcached:1.4.25'
restart: always
tags:
  - high-memory
```

crmrebs-prod-postgres:

```
environment:
  ...
image: 'rebs/postgresql:9.5-postgis'
restart: always
tags:
  - crmrebs-prod
  - ssd
```

crmrebs-prod-rabbitmq:

```
environment:
  ...
hostname: rabbitmq
image: 'rabbitmq:3.6.0-management'
restart: always
tags:
  - crmrebs-prod
  - ssd
```

crmrebs-prod-redis:

```
image: 'redis:3.0.7'
restart: always
tags:
  - crmrebs-prod
  - ssd
```

crmrebs-prod-web:

```
deployment_strategy: high_availability
environment:
  ...
image: 'rebs/crmrebs:master'
links:
  - 'crmrebs-prod-memcached:memcached'
  - 'crmrebs-prod-postgres:postgres'
  - 'crmrebs-prod-rabbitmq:rabbitmq'
  - 'crmrebs-prod-redis:redis'
  - 'telegraf.monitoring:telegraf'
mem_limit: 4096m
restart: always
sequential_deployment: true
target_num_containers: 5
```



# Docker Cloud - Stackfiles



crmrebs-prod-web-api:

```
command: ./docker-entrypoint-api.sh
deployment_strategy: high_availability
environment:
  ...
image: 'rebs/crmrebs:master'
links:
  - 'crmrebs-prod-memcached:memcached'
  - 'crmrebs-prod-postgres:postgres'
  - 'crmrebs-prod-rabbitmq:rabbitmq'
  - 'crmrebs-prod-redis:redis'
  - 'telegraf.monitoring:telegraf'
mem_limit: 2096m
restart: always
sequential_deployment: true
target_num_containers: 3
```

crmrebs-prod-worker-batch:

```
command: celery worker -A crmrebs -Q batch
deployment_strategy: high_availability
...
image: 'rebs/crmrebs:master'
links:
  - 'crmrebs-prod-memcached:memcached'
  - 'crmrebs-prod-postgres:postgres'
  - 'crmrebs-prod-rabbitmq:rabbitmq'
  - 'crmrebs-prod-redis:redis'
  - 'telegraf.monitoring:telegraf'
mem_limit: 4096m
restart: always
tags:
  - worker
target_num_containers: 2
```

crmrebs-prod-worker-priority:

```
command: celery worker -A crmrebs -Q priority
deployment_strategy: high_availability
...
image: 'rebs/crmrebs:master'
links:
  - 'crmrebs-prod-memcached:memcached'
  - 'crmrebs-prod-postgres:postgres'
  - 'crmrebs-prod-rabbitmq:rabbitmq'
  - 'crmrebs-prod-redis:redis'
  - 'telegraf.monitoring:telegraf'
mem_limit: 4096m
restart: always
tags:
  - worker
target_num_containers: 2
```



# Docker Cloud - Multi-host networking



- Overlay networking is now builtin in current Docker versions
  - ``overlay`` network driver
- Docker Cloud uses Weave
  - Before Docker implemented overlay networks
- All containers connected via a virtual network across all nodes
  - Zero configuration
  - All done seamlessly by Docker Cloud



# Docker Cloud - Service discovery



- DNS
  - Service discovery - using service name
  - Load balancing
- `dockercloud/haproxy`
  - Custom image tailored for Docker Cloud, uses the Docker Cloud API
  - Linked to your frontend Docker services
  - Uses linked service's environment variables for configuration



# Docker Cloud - Service discovery



master-haproxy:

```
image: 'dockercloud/haproxy:1.0.1'
deployment_strategy: high_availability
links:
  - 'crmrebs-prod-web.crmrebs-prod:crmrebs-prod-web'
  - 'crmrebs-prod-web-api.crmrebs-prod:crmrebs-prod-web-api'
  - 'jenkins.jenkins:jenkins'
  ...
ports:
  - '80:80'
  - '443:443'
restart: always
roles:
  - global
tags:
  - frontend
target_num_containers: 2
```

crmrebs-prod-web:

```
image: 'rebs/crmrebs:master'
...
environment:
  - 'HEALTH_CHECK=check inter 2000 rise 3 fall 2'
  - 'HTTP_CHECK=OPTIONS /login/ HTTP/1.1\r\nHost:\ demo.crmrebs.com'
  - 'VIRTUAL_HOST=*crmrebs.com,www.crmrebs.ro'
```

crmrebs-prod-web-api:

```
image: 'rebs/crmrebs:master'
...
environment:
  - VIRTUAL_HOST=api.crmrebs.com
```





# Docker Cloud - Deploying



crmrebs-prod-deploy:

command: `./docker-entrypoint-deploy.sh`

environment:

...

image: `'rebs/crmrebs:master'`

links:

- `'crmrebs-prod-memcached:memcached'`
- `'crmrebs-prod-postgres:postgres'`
- `'crmrebs-prod-rabbitmq:rabbitmq'`
- `'crmrebs-prod-redis:redis'`
- `'telegraf.monitoring:telegraf'`

crmrebs-prod-web:

deployment\_strategy: `high_availability`

sequential\_deployment: `true`

target\_num\_containers: `5`

...

```
# ./docker-entrypoint-deploy.sh
```

```
#!/bin/bash -e
```

```
echo "Running migrations..."
```

```
./manage.py migrate --noinput
```

```
echo "Uploading static files and compressing/minimizing static assets..."
```

```
./manage.py collectstatic --noinput && ./manage.py compress -e html -e js  
--force --engine jinja2 --engine django
```

```
echo "Loading permissions..."
```


```
./manage.py loaddata crmrebs/fixtures/permissions.json
```

```
echo "Loading tags..."
```

```
./manage.py loaddata crmrebs/fixtures/tags.json
```




# CI & CD with Jenkins & Docker



- Testing jobs
  - Runs tests inside a Docker container, same image used in prod environment
- Build job
  - git pull
  - docker build
  - docker push
- Deploy job - uses Docker Cloud API
  - Update Docker Cloud stack
  - Run one-off deploy container
  - Restart web & worker containers




# CI & CD with Jenkins & Docker



- Running Jenkins as a Docker service
  - Able to scale Jenkins slaves using Docker containers
  - Master and Slave Docker images
- Using Docker inside Docker containers - Dockerception
  - We do NOT use Docker-in-Docker
  - Link Docker socket and binary to the host's Docker daemon
  - Makes use of the host's Docker cache
  - Problem: mounting a directory as a data volume from within Jenkins will actually mount it on the host



# CI & CD with Jenkins & Docker



jenkins:

```
image: 'rebs/jenkins:latest'
```

```
expose:
```

- '8080'
- '50000'

```
links:
```

- jenkins-slave

```
restart: always
```

```
volumes:
```

- '/usr/bin/docker:/usr/bin/docker:ro'
- '/var/run/docker.sock:/var/run/docker.sock:ro'

jenkins-slave:

```
image: 'rebs/jenkins-slave:latest'
```

```
deployment_strategy: high_availability
```

```
target_num_containers: 2
```

```
volumes:
```

- '/usr/bin/docker:/usr/bin/docker:ro'
- '/var/run/docker.sock:/var/run/docker.sock:ro'



# Running multiple stacks

- Duplicate a Stackfile
  - Easily create as many isolated environments
- Stack configuration as code - custom deployer using Docker Cloud API

crmrebs-staging:

```
base: crmrebs
tag: develop
base_domain: staging-crmrebs.com
```

crmrebs-myfeature1:

```
base: crmrebs
tag: myfeature1
base_domain: myfeature1-crmrebs.com
```

crmrebs-myfeature2:

```
base: crmrebs
tag: myfeature2
base_domain: myfeature2-crmrebs.com
```



# Monitoring



- Using Docker to monitor Docker
- InfluxDB, Telegraf, Grafana stack
- Master Telegraf service which aggregates everything and writes to InfluxDB
- Run a Host Telegraf service on every host
  - Monitors host health
  - Monitors information about all running Docker containers
- For extra monitoring, point services to InfluxDB / Master Telegraf container
  - statsd style monitoring



# Monitoring



```
telegraf:
  image: 'rebs/telegraf:latest'
  links:
    - 'crmrebs-prod-memcached.crmrebs-prod:crmrebs-prod-memcached'
    - 'crmrebs-prod-postgres.crmrebs-prod:crmrebs-prod-postgres'
    - 'crmrebs-prod-rabbitmq.crmrebs-prod:crmrebs-prod-rabbitmq'
    - 'crmrebs-prod-redis.crmrebs-prod:crmrebs-prod-redis'
    - influxdb
    ...
  restart: always
```

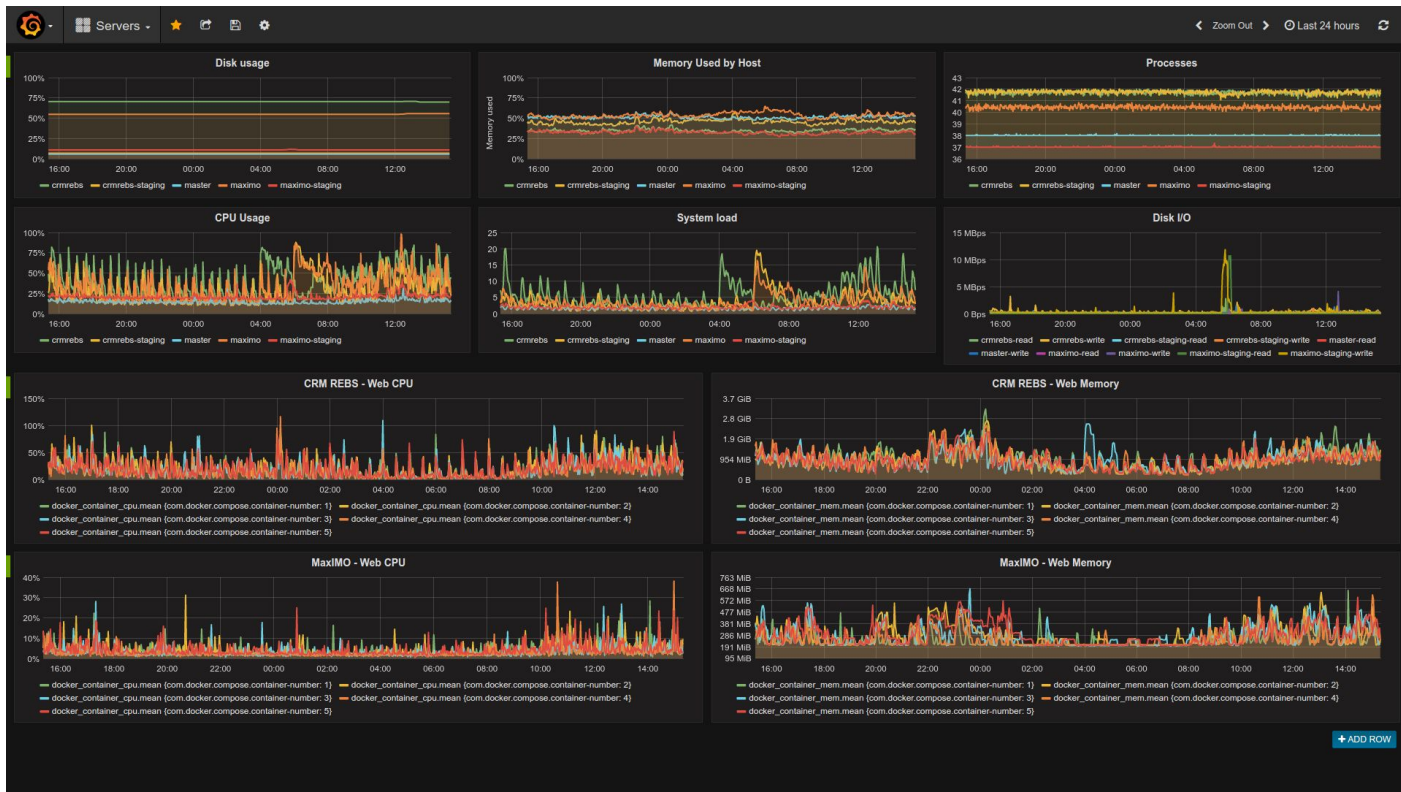
```
telegraf-host:
  image: 'rebs/telegraf-host:latest'
  deployment_strategy: every_node
  privileged: true
  restart: always
  volumes:
    - '/:/rootfs:ro'
    - '/proc:/mnt/host_proc:ro'
    - '/sys:/sys:ro'
    - '/var/lib/docker:/var/lib/docker:ro'
    - '/etc/hostname:/mnt/hostname:ro'
    - '/var/run/docker.sock:/var/run/docker.sock:ro'
```

---

```
crmrebs-prod-web:
  ...
  links:
    - 'telegraf.monitoring:telegraf'
    ...
```



# Monitoring







# Plans for the future



- Docker Swarm
  - Introduced in Docker v1.12
  - Orchestration
  - Built-in overlay network
  - Plugins
    - Network drivers - overlay
    - Volume plugins - Flocker, convoy



# Questions?

Email: [felix@crmrebs.com](mailto:felix@crmrebs.com)

Slides: <http://bit.ly/rebs-docker>



# REBS

Real Estate Business Solutions